

# **RV UNIVERSITY, BENGALURU**

## **SCHOOL OF COMPUTER SCIENCE AND ENGINEERING**



### **A Project Report On Analysis of Clash of Clans player communities**

B.SC(Honors)

In

School of Computer Science and Engineering

Submitted By

Team Member 01: 1RVU22BSC069 \_ PRABHAS BHAT

Team Member 02: 1RVU22BSC072 \_ PRASANNA G

Team Member 03: 1RVU22BSC052 \_ MERSHIKA U

Team Member 04: 1RVU22BSC021 \_ JYOSTHNA

**Big Data Analytics  
Under the Guidance of  
Vinod Kumar Raju  
Assistant Professor  
School of CSE  
RV University, Bengaluru-  
560059 2024-2025**

# **RV UNIVERSITY, BENGALURU-59**

## **SCHOOL OF COMPUTER SCIENCE AND ENGINEERING**



## **CERTIFICATE**

**Certified that the project work titled “Analysis of Clash of Clans player communities” is carried out by PRABHAS BHAT (1RVU22BSC069), PRASANNA G (1RVU22BSC072), MERSHIKA U (1RVU22BSC052), JYOSTHNA (1RVU22BSC021) RV University, Bengaluru, B.sc (Hons) in the School of Computer Science and Engineering of the RV University, Bengaluru during the year 2025-2026. It is certified that all corrections/suggestions indicated for the Internal Assessment have been incorporated in the project report. The Project report has been approved as it satisfies the academic requirements in respect of project work prescribed by the institution.**

### **Signature of Guide**

### **External Viva:**

**Name of Examiners**

**Signature with Date**

**1**

**2**

**RV UNIVERSITY, BENGALURU-59**

**SCHOOL OF COMPUTER SCIENCE AND ENGINEERING**

**DECLARATION**

We, Prabhas Bhat, Prasanna G, Mershika U, Jyosthna students of second semester BSc(Hons), SoCSE, RV University, Bengaluru, hereby declare that the project titled '**Analysis of Clash of Clans player communities**' has been carried out by us and submitted in partial fulfillment of **Bachelor of science(Hons) in School of Computer Science and Engineering** during the year 2025-26.

Further, we declare that the content of the report has not been submitted previously by anybody or to any other university.

We also declare that any Intellectual Property Rights generated out of this project carried out at RV University will be the property of RV University, Bengaluru, and we will be one of the authors of the same.

Place: Bengaluru

Date:

Name	Signature
1. PRABHAS BHAT (1RVU22BSC069)	
2. PRASANNA G (1RVU22BSC072)	
3. MERSHIKA U (1RVU22BSC052)	
4. JYOSTHNA (1RVU22BSC021)	

## **ACKNOWLEDGEMENT**

It is a great pleasure for us to acknowledge the assistance and support of many individuals who have been responsible for the successful completion of this project.

First, we take this opportunity to express our sincere gratitude to the School of Computer Science and Engineering, RV University, for providing us with a great opportunity to pursue our bachelor's degree in this institution.

A special thanks to our Program Director, **Dr. Sudhakar, and Dean - Dr.Shobha G**, for their continuous support and providing the necessary facilities with guidance to carry out mini project work.

We would like to thank our guide, Prof, Vinod Kumar Raju **Assistant Professor**, School of Computer Science and Engineering, RV University, for sparing his/her valuable time to extend help in every step of our project work, which paved the way for smooth progress and fruitful culmination of the project.

We are also grateful to our family and friends who provided us with every requirement throughout the course.

We would like to thank one and all who directly or indirectly helped us in the Project work.

*Signature of Student*

USN:

Name:

## Abstract

This project is a case study on a Clash of Clans dataset to analyse the distribution of clan types, and behaviour of each ranked league. We compared execution times for a single-node versus multi-node system, using a multi-node Apache Spark configuration on Google Cloud Platform with a single-node Apache Hive setup on a virtual machine. After running the same aggregation tasks on both systems for a considerable amount of clan data of 3.56 million rows, we came to the conclusion that a multi-node Spark is far more efficient than a single-node Hive setup, at least for such gaming big data analytics.

This study also highlights the importance of distributed computing frameworks in handling large-scale gaming datasets efficiently. By leveraging parallel processing and in-memory computation, Apache Spark demonstrates significant performance improvements over traditional single-node systems. The findings of this case study can serve as a reference for selecting suitable big data technologies for analytics-intensive applications in the gaming domain.

# Table of Contents:

	Page No
Abstract	iv
List of Tables	v
List of Figures	vi

<b>Page Title</b>	<b>Page No</b>
<b>Chapter 1: Introduction</b>	8
1.1. General Introduction	8
1.2. Literature Survey	8
1.3 Problem Statement / Objectives	8
<b>Chapter 2: System Design</b>	.9
2.1 Architectural Diagram	9
2.2 ER Modelling / ER Schema Diagram	9
<b>Chapter 3: Software Requirements</b>	9
3.1. Functional Requirements	9
3.2. Non-Functional Requirements	9
3.3. Hardware Requirements	9
3.4. Software Requirements	9
3.5. Summary	9
<b>Chapter 4: Implementation and Testing</b>	10
4.1. Modules	10
4.1.1. Descriptions of Modules	10
4.1.2. Aggregation Module	10
4.2 Testing	10
<b>Chapter 5: Results and Discussion</b>	10

5.1. Clan Type Distribution Results	10
5.2. League-wise Aggregation Results	10
5.3. Performance Comparison	10
5.4. Summary	10
<b>Chapter 6: Conclusion and Future Work</b>	12
<b>References [ APA format ]</b>	
<b>Appendices</b>	
<b>Appendix 1: Screenshots</b>	14
<b>Appendix 2: Source code</b>	20
<b>Appendix 3: Plagiarism details</b>	22

### **List of Tables**

<b>Table No</b>	<b>Table Name</b>	<b>Page No</b>
1.	Clan type aggregation	11
2.	Clan League aggregation	11

### **List of Figures**

<b>Figure No</b>	<b>Figure Name</b>	<b>Page No</b>
1.	Bar chart of execution times	12

# **1. Introduction**

## **1.1 General Introduction**

The recent increase in online multiplayer games has led to the generation of vast amounts of data concerning user behaviour, engagement, and interactions. Clash of Clans is one such globally recognized mobile game with millions of user bases divided into clans. The clans are diverse in terms of structure, accessibility, and level of competition, providing a superb platform to analyze large-scale interactions within a community via big data solutions.

In this particular project, big data analytics is used to analyze the data involving clans in Clash of Clans. The use of distributed computing platforms helps in gaining insights into types, leagues, and overall engagement.

## **1.2 Literature survey**

The relevance of big data platforms such as Apache Spark and Hive has been illustrated in a couple of research papers to efficiently process large chunks of data. Though Spark is known for its in-memory computation and parallelization properties, which make it apt for handling complex tasks, Hive is known for supporting SQL-type query processing on large chunks of data, but it is mostly limited by the disk execution environment when used on a single-node environment. The support research on gaming analytics focuses on developing scalable infrastructure that supports real-time analysis of game-related metrics such as player engagement, matchmaking, and competition balance.

## **1.3 Problem Statement & Objectives**

### **Problem Statement:**

The traditional single node data processing systems pose some challenges in terms of performance when dealing with massive gaming data. It is essential to assess whether the use of distributed computing platforms gives a significant advantage in such analytics tasks.

### **Objectives:**

- The analysis of the distribution of types in the Clash of Clans community.
- To analyze clan-related information in a league, including points, levels, and membership.
- The aim is to compare the performance of PySpark (Multi-node cluster using Google Cloud) with Hive (Single-node Cloudera VM).

## **2.System Design**

### **2.1 Architectural Diagram**

The architecture of the system comprises two parallel environments for processing:

- Spark on GCP: A multi-node cluster supporting Google Cloud Storage for storing the data with Spark for handling.

#### **Multi-node cluster configuration created on GCP:**

##### **Cluster specifications:**

Master node : n1 standard-4 (4 vCPUs, 15 GB memory), 100 GB harddisk

Worker Node: n1 standard 2 (2 vCPUs, 7.5 GB memory)

##### **Supports Jupyter Lab with PySpark**

- Hive on CloudEra VM: Single-node virtual machine with Apache Hive for SQL-based analytics.

##### **CloudEra VM specifications:**

Processor 1, Memory 5GB ,Hard Disk 64 GB

In summary, both environments ingest the same CSV file and run the same aggregation tasks.

### **2.2 Data Model / Schema Description**

The dataset consists of 27 attributes, including clan tag, clan name, clan type, clan war league, clan points, clan level, number of members, and war statistics. Aggregations are primarily performed on clan\_type and clan\_war\_league fields.

## **3.Software Requirements**

### **3.1 Functional Requirements**

- Load large CSV datasets into Spark and Hive environments.
- Perform aggregation and grouping operations on clan data.
- Measure and record execution time for each operation.
- Display and store analytical results.

### **3.2 Non-Functional Requirements**

- High performance and scalability.
- Fault tolerance in distributed execution.
- Accuracy and consistency of analytical results.

### **3.3 Hardware Requirements**

- Multi-node cluster on Google Cloud Platform.
- Virtual Machine with sufficient CPU, RAM, and storage.

### **3.4 Software Requirements**

- Apache Spark
- Apache Hive
- Python (PySpark)
- Google Cloud Platform (GCP)
- Cloudera VM with HIVE, Hadoop installed

### **3.5 Summary**

This chapter outlined the functional and non-functional requirements necessary to support large-scale data analytics using distributed and single-node systems.

## **4.Implementation and Testing**

### **4.1 Modules**

#### **4.1.1 Data Ingestion Module**

Loads the Clash of Clans CSV dataset from cloud storage into Spark DataFrames and Hive tables.

#### **4.1.2 Aggregation Module**

Performs:

- Simple aggregation on clan\_type.
- Complex aggregation on clan\_war\_league including averages of points, levels, and members.

#### **4.1.3 Performance Measurement Module**

Measures execution time using Python's time library to compare Spark and Hive performance.

### **4.2 Testing**

- Validation of row counts across both systems.
- Cross-verification of aggregation results.
- Performance testing using execution time comparisons.

## 5. Results and Discussion

### 5.1 Clan Type Distribution Results:

Clan Type	Count
open	2,569,771
inviteOnly	820,839
closed	169,133

Table-1: Clan type aggregation

Most clans are open, indicating a casual and inclusive community structure.

### 5.2 League-wise Aggregation Results

Unranked clans dominate the dataset, while higher leagues such as Champion and Master show fewer but more competitive clans with higher points, levels, and member counts.

League	Num Clans	Avg Clan Points	Avg Clan Level	Avg Num Members
Bronze League I	29,300	4124.03	3.38	11.50
Bronze League II	9,965	3912.77	3.13	12.52
Bronze League III	3,998	3749.62	2.94	13.54
Champion League I	556	22192.79	20.41	20.19
Champion League II	676	27267.77	20.01	24.20
Champion League III	1,480	30826.17	20.49	27.68
Crystal League I	21,187	31010.17	17.56	35.71
Crystal League II	22,230	25932.40	15.10	32.28
Crystal League III	26,156	21514.61	12.81	29.08
Gold League I	34,396	16801.08	10.60	24.33
Gold League II	49,293	12626.05	8.25	19.62
Gold League III	70,027	9165.39	6.17	15.29
Master League I	2,963	35091.31	21.34	33.04
Master League II	5,863	34971.71	20.60	34.95
Master League III	11,288	33625.49	19.39	35.84
Silver League I	81,990	6925.99	4.93	12.72
Silver League II	77,200	5530.49	4.25	11.45
Silver League III	75,017	4539.81	3.70	10.93
Unranked	3,036,158	1072.62	1.49	2.57

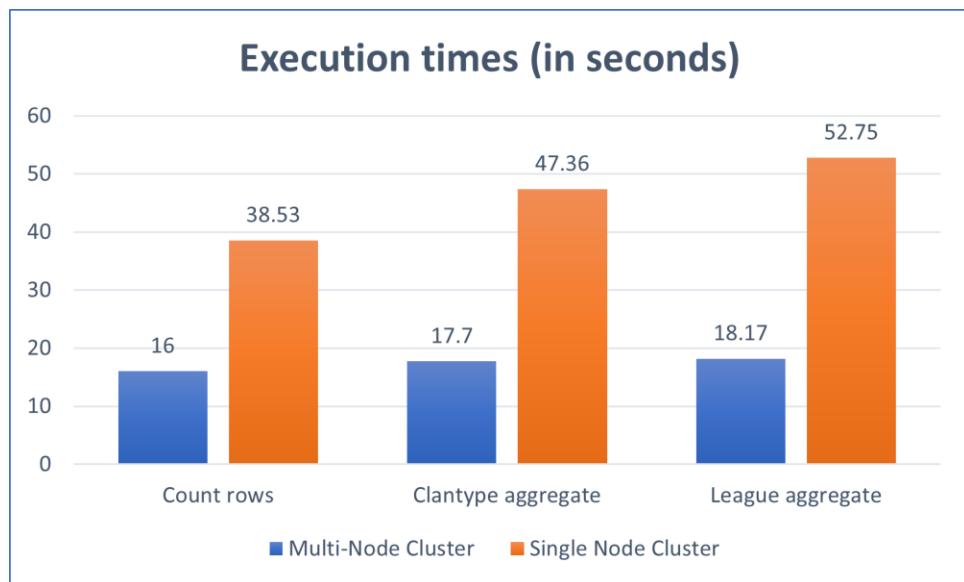
Table-2: Clan league aggregation

### 5.3 Performance Comparison

The Spark cluster consistently outperformed the single-node Hive environment:

- Up to 5×–10× faster execution times.
- Better scalability for large aggregations.

**Execution times comparison:**



**Figure-1:** Bar chart of execution times

### 5.4 Summary

**Business Problem perspective:** Open clans dominate the ecosystem, highlighting a preference for accessible and casual gameplay communities. In contrast, invite-only and closed clans represent more structured and controlled participation, indicating varied engagement strategies within the player base.

**Performance perspective:** The row count operation executed **5.5x faster** on the multi-node Spark cluster, while clan type aggregation, which involves a GROUP BY operation, showed a **4x** performance improvement on Spark. The league-wise aggregation, which includes multiple aggregate functions (COUNT, AVG) ran almost **3x faster** on multi-node Spark.

## 6. Conclusion and Future Work

### Conclusion

The results prove that Apache Spark running on a multi-node GCP cluster provided substantial performance benefits over a single-node Hive setup when dealing with big datasets. The Spark cluster benefits from **parallel execution, in-memory computation, and efficient task scheduling**, while the single-node Hive setup is limited by disk I/O and sequential processing.

### Future Work

- Including real-time streaming data.
- Applying machine learning models for clan behaviour prediction.
- Extending analysis to player-level engagement metrics.

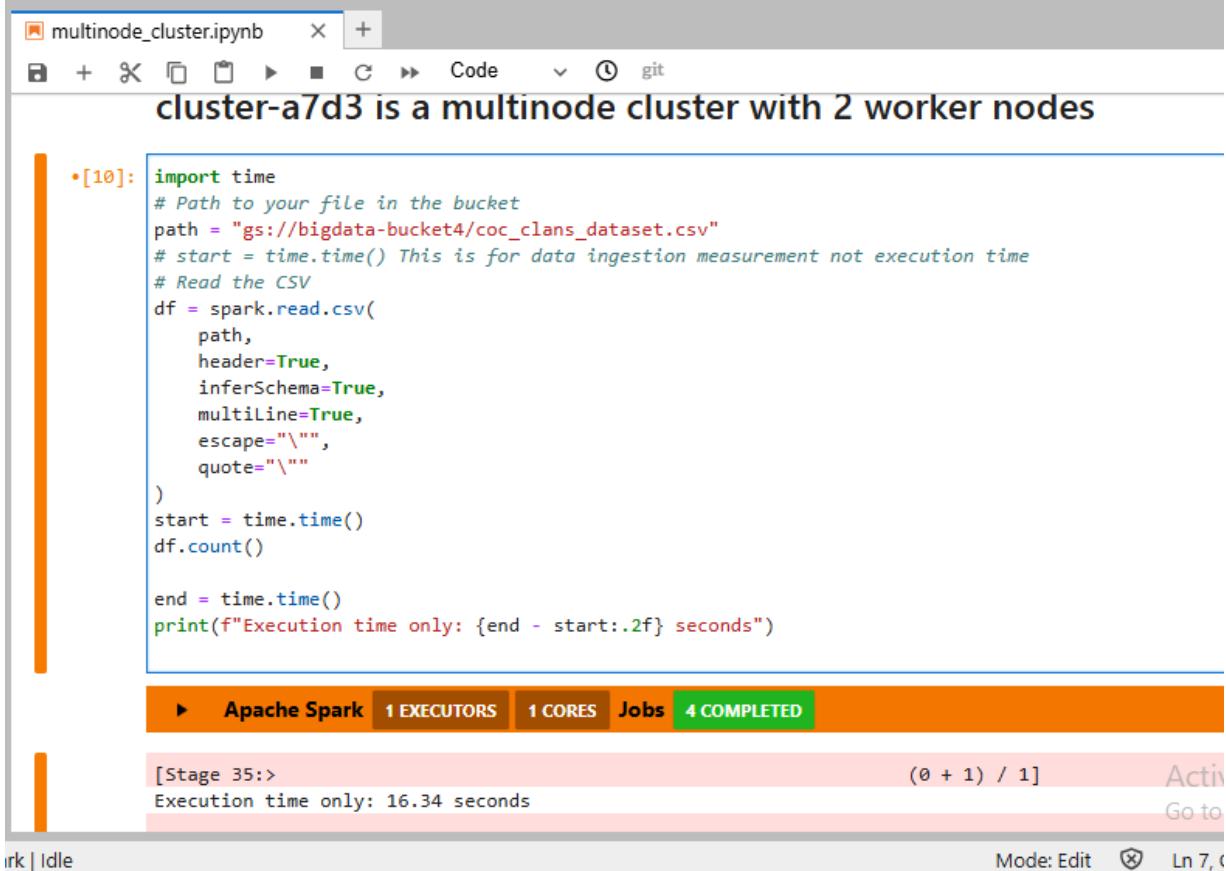
### References

- Zaharia, M. et al. (2016). *Apache Spark: A Unified Engine for Big Data Processing*. Communications of the ACM.
- Kaggle. (2023). *Clash of Clans Clans Dataset*.

## Appendices

Appendix 1: Screenshots of Multi-node GCP and CloudEra VM operations.

### 1)Multi-node row count operation



```
multinode_cluster.ipynb + Code git
cluster-a7d3 is a multinode cluster with 2 worker nodes

• [10]: import time
# Path to your file in the bucket
path = "gs://bigdata-bucket4/coc_clans_dataset.csv"
# start = time.time() This is for data ingestion measurement not execution time
# Read the CSV
df = spark.read.csv(
    path,
    header=True,
    inferSchema=True,
    multiLine=True,
    escape="\"",
    quote="\""
)
start = time.time()
df.count()

end = time.time()
print(f"Execution time only: {end - start:.2f} seconds")
```

► Apache Spark 1 EXECUTORS 1 CORES Jobs 4 COMPLETED

[Stage 35:> (0 + 1) / 1] Active Go to

Execution time only: 16.34 seconds

Work | Idle Mode: Edit Ln 7, C

## 2) Multi-Node clantype aggregation

```
[5]: import time
      start = time.time()
      clan_type_dist = (
          df.groupBy("clan_type")
            .count()
            .orderBy("count", ascending=False)
      )
      clan_type_dist.show()
      end = time.time()
      print(f"\nTime taken for clan_type aggregation: {end - start:.3f} seconds")
```

▶ Apache Spark 1 EXECUTORS 1 CORES Jobs 2 COMPLETED

```
[Stage 18:>
+-----+-----+
| clan_type| count|
+-----+-----+
|    open|2569771|
|inviteOnly| 820839|
|   closed| 169133|
+-----+-----+
```

```
Time taken for clan_type aggregation: 17.780 seconds
```

### 3) Multi-node Clan League Aggregation

```
[6]: from pyspark.sql import functions as F
import time
start = time.time()

league_stats = (
    df.groupBy("clan_war_league")
    .agg(
        F.count("*").alias("num_clans"),
        F.avg("clan_points").alias("avg_clan_points"),
        F.avg("clan_level").alias("avg_clan_level"),
        F.avg("num_members").alias("avg_num_members")
    )
    .orderBy("clan_war_league")
)
league_stats.show()
end = time.time()
print(f"\nTime taken for war league aggregation: {end - start:.2f} seconds")
```

▶ Apache Spark 1 EXECUTORS | 1 CORES | Jobs 2 COMPLETED

[Stage 21:> (0 + 1) / 1]				
clan_war_league	num_clans	avg_clan_points	avg_clan_level	avg_num_members
Bronze League I	29300	4124.027576791809	3.3819453924914678	11.501160409556315
Bronze League II	9965	3912.768188660311	3.1250376317109883	12.515203211239339
Bronze League III	3998	3749.6193096548272	2.94072036018009	13.536268134067033
Champion League I	556	22192.793165467625	20.408273381294965	20.190647482014388
Champion League II	676	27267.76923076923	20.014792899408285	24.20414201183432
Champion League III	1480	30826.16689189189	20.49256756756757	27.681756756756755
Crystal League I	21187	31010.17397460707	17.556426110350685	35.706376551659034
Crystal League II	22230	25932.39851551957	15.09676113360324	32.278542510121454
Crystal League III	26156	21514.61370240098	12.811209665086405	29.08338430952745
Gold League I	34396	16801.075735550647	10.599197581114083	24.333934178392838
Gold League II	49293	12626.054105045341	8.245044935386364	19.622400746556306
Gold League III	70027	9165.393548202836	6.166764248075742	15.286046810515943
Master League I	2963	35091.31488356396	21.337833277084037	33.03543705703679
Master League II	5863	34971.713627835576	20.5973051338905	34.95480129626471
Master League III	11288	33625.49317859674	19.39466690290574	35.837615166548545
Silver League I	81990	6925.987681424564	4.928113184534699	12.718063178436395
Silver League II	77200	5530.494352331606	4.247810880829015	11.4464896373057
Silver League III	75017	4539.810229681272	3.7039604356345897	10.93166882173374
Unranked	3036158	1072.6167959638465	1.4852761285809237	2.5657831377681926

Time taken for war league aggregation: 18.17 seconds

Activate W  
Go to Settings

#### 4)Hive Simple row Count

```
cloudera@quickstart:~/Desktop
File Edit View Search Terminal Help
Silver League III      75017    4539.810229681272      3.7039604356345897  1
0.93166882173374
Unranked          3036158 1072.6167959638465      1.4852761285809237      2.565783
1377681926
Time taken: 49.196 seconds, Fetched: 19 row(s)
hive> SELECT COUNT(*) FROM clash_data;
Query ID = cloudera_20251214120707_b7a4c0e3-08cb-4c35-83b6-b43c12a72c7f
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1765734899499_0019, Tracking URL = http://quickstart.cloudera
:8088/proxy/application_1765734899499_0019/
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1765734899499_0019
Hadoop job information for Stage-1: number of mappers: 2; number of reducers: 1
2025-12-14 12:07:36,593 Stage-1 map = 0%,  reduce = 0%
2025-12-14 12:07:54,485 Stage-1 map = 50%,  reduce = 0%, Cumulative CPU 3.02 sec
2025-12-14 12:07:55,545 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 7.16 se
c
```

```
cloudera@quickstart:~/Desktop
File Edit View Search Terminal Help
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1765734899499_0019, Tracking URL = http://quickstart.cloudera
:8088/proxy/application_1765734899499_0019/
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1765734899499_0019
Hadoop job information for Stage-1: number of mappers: 2; number of reducers: 1
2025-12-14 12:07:36,593 Stage-1 map = 0%,  reduce = 0%
2025-12-14 12:07:54,485 Stage-1 map = 50%,  reduce = 0%, Cumulative CPU 3.02 sec
2025-12-14 12:07:55,545 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 7.16 se
c
2025-12-14 12:08:04,083 Stage-1 map = 100%,  reduce = 100%, Cumulative CPU 8.44
sec
MapReduce Total cumulative CPU time: 8 seconds 440 msec
Ended Job = job_1765734899499_0019
MapReduce Jobs Launched:
Stage-Stage-1: Map: 2 Reduce: 1 Cumulative CPU: 8.44 sec HDFS Read: 4151408
43 HDFS Write: 8 SUCCESS
Total MapReduce CPU Time Spent: 8 seconds 440 msec
OK
3559743
Time taken: 38.538 seconds, Fetched: 1 row(s)
hive>
```

## 5)Hive clantype aggregation

```
cloudera@quickstart:~/Desktop
File Edit View Search Terminal Help
Stage-Stage-1: Map: 2 Reduce: 1 Cumulative CPU: 8.44 sec HDFS Read: 4151408
43 HDFS Write: 8 SUCCESS
Total MapReduce CPU Time Spent: 8 seconds 440 msec
OK
3559743
Time taken: 38.538 seconds, Fetched: 1 row(s)
hive> SELECT
>     clan_type,
>     COUNT(*) AS clan_count
>   FROM coc_clean
> GROUP BY clan_type
> ORDER BY clan_count DESC;
Query ID = cloudera_20251214121010_c4c15d0e-2037-4781-83e3-d645d58d0f77
Total jobs = 2
Launching Job 1 out of 2
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
```

```
cloudera@quickstart:~/Desktop
File Edit View Search Terminal Help
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1765734899499_0021, Tracking URL = http://quickstart.cloudera
:8088/proxy/application_1765734899499_0021/
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1765734899499_0021
Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1
2025-12-14 12:11:07,972 Stage-2 map = 0%, reduce = 0%
2025-12-14 12:11:14,414 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 0.8 sec
2025-12-14 12:11:21,803 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 1.99
sec
MapReduce Total cumulative CPU time: 1 seconds 990 msec
Ended Job = job_1765734899499_0021
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 4.83 sec HDFS Read: 1307716
67 HDFS Write: 182 SUCCESS
Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 1.99 sec HDFS Read: 5107 HD
FS Write: 45 SUCCESS
Total MapReduce CPU Time Spent: 6 seconds 820 msec
OK
open    2569771
inviteOnly    820839
closed   169133
Time taken: 47.369 seconds, Fetched: 3 row(s)
hive> ■
```

## 6) Hive Clan League aggregate

```
cloudera@quickstart:~/Desktop
File Edit View Search Terminal Help
closed 169133
Time taken: 47.369 seconds, Fetched: 3 row(s)
hive> SELECT
    >     clan_war_league,
    >     COUNT(*) AS num_clans,
    >     AVG(clan_points) AS avg_clan_points,
    >     AVG(clan_level) AS avg_clan_level,
    >     AVG(num_members) AS avg_num_members
    > FROM coc_clean
    > GROUP BY clan_war_league
    > ORDER BY clan_war_league;
Query ID = cloudera_20251214121313_02602609-1a0e-4b2b-9433-465c8c042e66
Total jobs = 2
Launching Job 1 out of 2
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducer.bytes.per.reducer.number</pre>
```

```
cloudera@quickstart:~/Desktop
File Edit View Search Terminal Help
Ended Job = job_1765734899499_0025
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 5.58 sec HDFS Read: 130773460 HDFS Write: 1255 SUCCESS
Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 1.94 sec HDFS Read: 7369 HDFS Write: 1477 SUCCESS
Total MapReduce CPU Time Spent: 7 seconds 520 msec
OK
Bronze League I 29300 4124.027576791809 3.3819453924914678 11.501160409556315
Bronze League II 9965 3912.768188660311 3.1250376317109883 12.515203211239339
Bronze League III 3998 3749.6193096548272 2.94072036018009 13.536268134067033
Champion League I 556 22192.793165467625 20.408273381294965 20.190647482014388
Champion League II 676 27267.76923076923 20.014792899408285 24.20414201183432
Champion League III 1480 30826.16689189189 20.49256756756757 27.681756756756755
Crystal League I 21187 31010.17397460707 17.556426110350685 35.706376551659034
Crystal League II 22230 25932.39851551957 15.09676113360324 32.278542510121454
Crystal League III 26156 21514.61370240098 12.811209665086405 29.08338430952745
Gold League I 34396 16801.075735550647 10.599197581114083 24.333934178392838
Gold League II 49293 12626.054105045341 8.245044935386364 19.622400746556306
Gold League III 70027 9165.393548202836 6.166764248075742 15.286046810515943
Master League I 2963 35091.31488356396 21.337833277084037 33.03543705703679
Master League II 5863 34971.713627835576 20.5973051338905 34.95480129626471
Master League III 11288 33625.49317859674 19.39466690290574 35.837615166548545
Silver League I 81990 6925.987681424564 4.928113184534699 12.718063178436395
Silver League II 77200 5530.494352331606 4.247810880829015 11.4464896373057
Silver League III 75017 4539.810229681272 3.7039604356345897 10.93166882173374
Unranked 3036158 1072.6167959638465 1.4852761285809237 2.5657831377681926
Time taken: 52.75 seconds, Fetched: 19 row(s)
hive> S
```

## **Appendix 2:**

### **GCP Source Code:**

#### **#Count rows**

```
import time
# Path to your file in the bucket
path = "gs://bigdata-bucket4/coc_clans_dataset.csv"
# start = time.time() This is for data ingestion measurement not execution time
# Read the CSV
df = spark.read.csv(
    path,
    header=True,
    inferSchema=True,
    multiLine=True,
    escape="\""",
    quote="\"""
)
start = time.time()
df.count()

end = time.time()
print(f"Execution time only: {end - start:.2f} seconds")
```

#### **#Clan type aggregation**

```
from pyspark.sql import functions as F
```

```
import time
start = time.time()
clan_type_dist = (
    df.groupBy("clan_type")
    .count()
    .orderBy("count", ascending=False)
)
clan_type_dist.show()
end = time.time()
print(f"\nTime taken for clan_type aggregation: {end - start:.3f} seconds")
```

#### **#Clan League aggregation**

```
from pyspark.sql import functions as F
```

```
import time
start = time.time()
```

```
league_stats = (
    df.groupBy("clan_war_league")
    .agg(
        F.count("*").alias("num_clans"),
        F.avg("clan_points").alias("avg_clan_points"),
        F.avg("clan_level").alias("avg_clan_level"),
        F.avg("num_members").alias("avg_num_members")
    )
    .orderBy("clan_war_league")
)
league_stats.show()
```

```
end = time.time()
print(f"\nTime taken for war league aggregation: {end - start:.2f} seconds")
```

### Hive Source Code:

#### Steps for Hive analysis:

Hadoop Filesystem: hadoop fs -put 'clash\_clean.csv'

#### Initialised hive in terminal

**Step-1:**Creating table and inserting data:

```
CREATE TABLE clash_data ( clan_tag STRING, clan_name STRING, clan_type STRING,
clan_location STRING, isFamilyFriendly STRING, clan_level INT, clan_points INT,
clan_builder_base_points INT, clan_versus_points INT, required_trophies INT, war_frequency
STRING, war_win_streak INT, war_wins INT, war_ties INT, war_losses INT, clan_war_league
STRING, num_members INT, required_builder_base_trophies INT, required_versus_trophies INT,
required_townhall_level INT, clan_capital_hall_level INT, clan_capital_points INT, capital_league
STRING, mean_member_level DOUBLE, mean_member_trophies DOUBLE ) ROW FORMAT
DELIMITED FIELDS TERMINATED BY ',' STORED AS TEXTFILE;
```

#### **STEP-2:** (Load csv data from hadoop into the table) :

```
LOAD DATA INPATH 'clash_clean_25.csv' INTO TABLE clash_data;
```

#### **STEP-3: (Verify data insertion)**

```
SELECT COUNT(*) FROM clash_data;
```

#### **STEP-4: Clantype aggregation:**

```
SELECT clan_type, COUNT(*) AS clan_count FROM coc_clean GROUP BY clan_type ORDER
BY clan_count DESC;
```

#### **STEP-5: Clan league aggregation:**

```
SELECT clan_war_league, COUNT(*) AS num_clans, AVG(clan_points) AS avg_clan_points,
AVG(clan_level) AS avg_clan_level, AVG(num_members) AS avg_num_members FROM
coc_clean GROUP BY clan_war_league ORDER BY clan_war_league;
```

### Appendix 3: Plagiarism Details



Plagiarism

8%

Partial Match

2%

Exact Match

6%

Unique

92%

#### Scan details

Total Words	Total Characters	Plagiarized Sentences	Unique Sentences
1025	7023	3.28	37.72 (92%)

#### Plagiarism Results: (4)

#1 2% Similar

<https://massedcompute.com/faq-answers/?ques tio...>

\* Virtual Machine with sufficient CPU, RAM, and storage.

#2 2% Similar

<https://massedcompute.com/faq-answers/?ques tio...>

\* Virtual Machine with sufficient CPU, RAM, and storage.

#3 2% Similar

<https://security.cbidigital.com/uploaded-files/e7671...>

\* Virtual Machine with sufficient CPU, RAM, and storage.

#4 2% Similar

<https://www.databricks.com/research/apache-spar...>

Apache Spark: A Unified Engine for Big Data Processing.